

PvA Position paper on use of AI for simplification of municipal communication.

FRANK VAN DER VELDE, Hogeschool van Amsterdam, Netherlands

This position paper addresses a policy question currently faced by Dutch municipalities, illustrated through the case of the Municipality of Oostzaan: should AI assisted language tools be structurally used to simplify municipal communication to B1 level in order to improve accessibility for residents, or does such reliance carry unacceptable long-term linguistic risks? While recent research demonstrates that large language models can produce B1-level texts [18, 24], the central issue is whether institutionalizing their use as routine rewriting instruments aligns with broader public responsibilities concerning language, comprehension, and civic resilience.

One policy position emphasises efficiency and accessibility. From this view, AI assisted rewriting supports compliance with national B1 guidelines promoted by CommunicatieRijk [9, 32], reduces staff workload, and improves comprehension for residents with limited literacy or language proficiency, concerns highlighted by organisations such as Stichting Lezen & Schrijven and Pharos [31, 36]. A contrasting position cautions that structural reliance on AI-driven simplification risks reinforcing linguistic homogenization, narrowing civic language, and shifting responsibility away from addressing deeper educational and communicative challenges, concerns grounded in critiques of CEFR-based simplification and emerging research on AI-generated language [5, 38].

This paper evaluates multiple policy options and argues against structural AI authorship or routine automated rewriting of official texts. Instead, it defends a limited, diagnostic role for AI that supports human judgment and institutional language capacity, based on considerations of municipal responsibility, linguistic resilience, and maintenance of human writing skill.

Additional Key Words and Phrases: AI text simplification, municipal communication, Dutch language, B1 language level, linguistic homogenization, large language models, language policy

ACM Reference Format:

Frank van der Velde. 2018. PvA Position paper on use of AI for simplification of municipal communication.. *Proc. ACM Meas. Anal. Comput. Syst.* 37, 4, Article 111 (August 2018), 14 pages. <https://doi.org/XXXXXXX.XXXXXXX>

Author's Contact Information: Frank van der Velde, frank.van.der.velde@hva.nl, Hogeschool van Amsterdam, Amsterdam, Netherlands.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM 2476-1249/2018/8-ART111

<https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

This paper addresses a recent policy question faced by Dutch municipalities: how should AI-based language tools be integrated into municipal communication without undermining public responsibility for civic language? Using the Municipality of Oostzaan as an illustrative case, it evaluates several policy options and defends a clear position. It argues against structural AI authorship and routine automated rewriting to B1 level, and instead supports a restricted, diagnostic use of AI that strengthens human judgment and institutional capacity.

Dutch municipalities are responsible for communicating clearly, accurately, and accessibly with residents about permits, benefits, taxes, and enforcement decisions. Over the past decade, this responsibility has increasingly been realised through the national CEFR B1-targeted language policy. National frameworks developed by CommunicatieRijk promote B1 as the default level for public communication [32], reinforced by broader commitments to inclusive communication [9] and by guidance from organisations such as Pharos and Stichting Lezen & Schrijven, which argue that simplified language benefits a wide range of readers [31, 36].

In practice, meeting this benchmark is demanding. Municipal texts often involve abstract rules, legal conditions, and procedural obligations that resist straightforward simplification. Employees differ in linguistic expertise, and rewriting texts to comply with B1 guidelines requires time and specialised skill. Empirical research shows that misunderstandings persist even where simplification efforts are made. Corsius and colleagues demonstrate in the Monitor Begrijpelijkheid overheidsteksten that citizens across educational levels struggle with bureaucratic phrasing and institutional logic [10]. Similar findings are reported by the Kenniscentrum Begrijpelijke Taal, which emphasises that accessibility problems are structurally perceived in administrative communication [22].

Against this backdrop, AI-based writing and rewriting tools are increasingly presented as practical solutions. Large language models can rapidly simplify texts, generate alternative formulations, and check alignment with writing guidelines. Initiatives within national government, including at the Belastingdienst and the Toeslagen system, illustrate how AI assisted simplification is already being explored [6, 7]. For municipalities facing limited resources, such tools promise efficiency, standardisation, and visible compliance with accessibility norms.

This promise introduces a policy tension that cannot be resolved through technical evaluation alone. Evidence of declining literacy in the Netherlands, reported by the Algemene Rekenkamer and the Inspectie van het Onderwijs, indicates that language difficulties are part of a broader societal challenge [1, 19]. At the same time, critiques of the CEFR framework question the suitability of B1 as a proxy for text complexity outside language assessment contexts [38]. Emerging research on AI-generated language further suggests that routine automated rewriting may contribute to linguistic standardisation over time [5].

For municipalities, the central question is therefore not whether AI can be useful, but how its use should be governed. This position paper moves beyond feasibility and efficiency considerations and instead provides decision a grounded guidance based on responsibility. It identifies and compares distinct policy positions on the use of AI assisted language tools in municipal communication, situating each position in relevant academic literature, policy documents, and applied AI research. Rather than testing hypotheses or reporting empirical results, the analysis evaluates trade offs between accessibility, institutional responsibility, and long term linguistic effects, with the aim of supporting informed decision making by municipal professionals through clarification of underlying assumptions, risks, and consequences. On this basis, the paper explicitly argues for a restrictive governance position, namely that AI should be used exclusively for diagnostic and learning oriented purposes, rather than for structural or automated simplification of municipal communication.

2 Base understanding of LLM's

This section explains why the technical properties of large language models matter for municipal decision making. A basic understanding of how these systems generate text helps clarify why certain linguistic effects arise when they are used at institutional scale.

Large language models are trained on extensive corpora of existing texts drawn from public, professional, and institutional sources. Through this process, they learn statistical regularities in language use, including which formulations commonly occur in specific contexts. When generating text, these models produce sequences of words that are statistically likely given patterns observed in their training data, as described in applied and technical research on AI writing systems [2].

This probabilistic mechanism has predictable consequences. Because dominant formulations appear more frequently in training data, AI systems tend to reproduce standardised expressions rather than rare or context-specific ones. Empirical studies show that repeated use of such systems leads to convergence toward common lexical and syntactic patterns, even when users retain formal control over final output [5]. When AI tools are used routinely for similar tasks, such as rewriting administrative correspondence, this tendency is amplified.

For municipal communication, these effects are not neutral. Official language shapes how rights, obligations, and procedures are understood by citizens. When AI-generated phrasing becomes routine, it can gradually set implicit norms for how a municipality communicates, even without explicit policy intent. Research on AI assisted writing suggests that exposure to homogenized language influences how users structure their own texts and allocate attention [23]. At institutional scale, such shifts accumulate through repetition rather than through any single decision.




Fig. 1. Illustration of how repeated mediation by the same large language model can reduce variation in how people express ideas, frame perspectives, and structure reasoning. When diverse users rely on a shared AI system for rewriting or polishing text, outputs tend to converge toward standardised formulations over time. Reproduced from Sourati et al. [35].

These dynamics reflect structural properties of probabilistic text generation applied repeatedly within an organisation. For municipalities, the central question is therefore not whether large language models function correctly, but how their routine use interacts with public responsibility for civic language, clarity, and clear and targeted communication.

3 Existing literature and argumentative groundwork

This section surveys the literature that underpins the paper's argument. It first documents the structural accessibility problems in municipal communication and the resulting case for simplification. It then reviews research

questioning the adequacy of B1 and CEFR-based approaches as policy tools, followed by studies showing that AI-generated language tends toward homogenization. The section concludes with historical and comparative work illustrating how institutional language choices produce durable societal effects.

3.1 Evidence supporting the need for simplification

A substantial body of literature supports the claim that Dutch municipalities face a structural accessibility problem. Reports by the Algemene Rekenkamer show that approximately 2.5 million adults in the Netherlands experience difficulties with reading and numeracy, directly limiting their ability to understand official correspondence and participate fully in civic life [1]. Communication barriers are therefore not isolated incidents but a systemic issue with societal consequences.

International and national assessments reinforce this picture. OECD PISA results document a sustained decline in reading performance among Dutch pupils, with an increasing share failing to meet baseline proficiency levels, a trend that has intensified relative to comparable countries [27, 29]. Broader system-level analyses, including Education at a Glance, highlight persistent inequalities linked to educational background and literacy [28]. National oversight bodies report similar patterns. The Inspectie van het Onderwijs shows that many students do not reach intended reference levels in reading and that literacy outcomes remain under pressure across educational tracks, with widening differences between schools [19–21]. Together, these sources indicate that literacy challenges are structural rather than temporary.

Accessibility problems also extend beyond traditionally defined low-literate groups. Research commissioned by the Nederlandse Taalunie demonstrates that even higher-educated citizens struggle with bureaucratic phrasing, legal abstraction, and institutional logic. The Monitor Begrijpelijkheid overheidsstukken shows that misunderstandings persist across domains such as healthcare and financial communication despite longstanding plain-language policies [10]. This suggests that communicative failures persist not just because of individual deficits, but because existing plain-language policies have also been insufficient to address the structural features of administrative language.

Policy-oriented organisations therefore argue that simplification benefits a broad audience. CommunicatieRijk promotes B1-level writing as a pragmatic default that improves comprehension and reduces error for most readers [32]. Pharos similarly emphasises that clear language supports people under stress, NT2 speakers, and individuals with limited language skills while remaining acceptable to higher-educated readers [31]. Stichting Lezen & Schrijven adds that even B1 remains too complex for some groups, underscoring the scale of the accessibility challenge [12, 36]. Viewed together, this literature establishes strong ethical and practical reasons for municipalities to prioritise accessibility.

3.2 Evidence questioning B1 and CEFR as sufficient solutions

A second body of literature complicates the assumption that targeting B1 automatically resolves accessibility problems. The Common European Framework of Reference was developed as a descriptive scale for learner proficiency in language education, not as a tool for diagnosing the intrinsic complexity of administrative texts [11]. Its levels support curriculum alignment and assessment rather than real-world text comprehensibility.

Scholarly critiques highlight the limitations of extending CEFR beyond this scope. Weir argues that CEFR descriptors abstract away from contextual, cognitive, and task-related variables that strongly influence comprehension, particularly in unfamiliar or high-stakes situations [38]. Assigning a single level such as B1 to a text therefore risks oversimplifying the interaction between language form, content, and reader background.

Empirical research reinforces this concern. Morales shows that grammatical and structural complexity varies substantially within CEFR bands and that texts classified at the same level can impose markedly different processing demands [25]. Siripol and colleagues further demonstrate that automated CEFR analyzers frequently

assign inconsistent levels to identical texts, with results sensitive to genre and surface features [34]. These findings raise doubts about the reliability of CEFR-based classification for policy enforcement or quality control.

Research on government communication supports this critique from another angle. The Monitor Begrijpelijkheid overheidsstukken shows that misunderstandings often arise from institutional framing, implicit assumptions, and legal logic rather than from vocabulary or sentence length alone [10]. Guidance from the Kenniscentrum Begrijpelijke Taal similarly emphasises that clarity depends on information structure and contextualisation, not merely linguistic simplification [22]. Together, this literature calls into question if the B1 CEFR level functions as a sufficient or reliable solution to accessibility challenges.

3.3 Evidence that AI-generated language leads to homogenization

Prior research has shown that while large language models can improve surface-level correctness and readability, their outputs tend to converge toward statistically dominant patterns of expression. Corpus studies by Fredrick and Craven demonstrate that texts generated with systems such as ChatGPT exhibit reduced lexical and syntactic variation compared to human-authored writing, including writing by language learners [15]. Homogenization therefore appears as a characteristic of the technology rather than a by-product of low proficiency.

Applied studies in educational settings report similar effects. Gayed and colleagues find that AI assisted writing improves grammatical accuracy and cohesion but encourages formulaic phrasing and reduced engagement with content [16]. Farhan's content analysis likewise shows that overreliance on AI tools weakens independent editing skills and expressive diversity despite improvements in surface quality [14]. These findings indicate that AI assistance reshapes writing practices rather than merely correcting errors.

Research on text simplification reinforces this pattern. Studies by Ardila and Kakarmath and by Guidroz and colleagues show that LLM-based simplification improves comprehension and reduces cognitive load [3, 18]. However, these gains are achieved by systematically standardising linguistic features. Mo and Hu show that high-performing simplification systems optimise for consistency and predictability across reading levels [24], aligning closely with homogenization dynamics.

More explicitly critical work frames these tendencies as a societal risk. Bauer argues that widespread adoption of ChatGPT-style systems increases convergence in academic and professional writing [5]. Sourati and colleagues extend this concern by suggesting that reliance on shared language models may reduce cognitive and linguistic diversity at scale [35]. Institutional deployment magnifies these effects. Experimental findings discussed by the MIT Media Lab show reduced neural engagement and increased textual similarity when users write with LLM assistance [23]. Embedded in municipal workflows, such systems risk normalising a narrow administrative register.

3.4 Historical and comparative cases of institutional language shaping

Historical and comparative research shows that institutional language choices can have long-term and often unintended consequences. Studies of language policy demonstrate that when public institutions systematically privilege particular forms of language, these choices reshape linguistic hierarchies, everyday usage, and perceptions of legitimacy over time.

A frequently cited case is Singapore. Bokhorst-Heng and Silver analyse how state language policies promoting English and Mandarin reshaped linguistic practices within the Chinese community, even as official discourse framed these policies as pragmatic and economically necessary [8]. Their analysis shows that once institutional preferences are embedded in education and administration, they can accelerate language shift and marginalise alternative varieties despite formal commitments to multilingualism [13, 37]. The relevance of this case lies not in its multilingual setting, but in its demonstration that language policy effects accumulate through routine institutional use rather than explicit coercion.

Similar dynamics are documented in other national contexts. Ng’s analysis of mother-tongue education in Singapore shows how English gradually displaced other languages in practice, despite policy rhetoric emphasising balance and additive bilingualism [26]. Gonzalez’s account of language planning in the Philippines likewise illustrates how administrative and educational priorities narrowed the functional domains of local languages over time, often in ways that proved difficult to reverse once institutional norms were established [17]. Together, these studies highlight a recurring pattern in which pragmatic institutional language choices produce durable effects on linguistic diversity and competence.

More recent work extends these insights to contemporary debates about language, power, and technology. A decolonial review of English language teaching in the Philippines shows how global language infrastructures reinforce dominant norms even when reforms explicitly aim to promote linguistic justice [4]. Although this literature does not address AI directly, it underscores how technologies and policies that appear neutral often reproduce existing hierarchies unless actively governed.

From a broader sociocultural perspective, Roche’s analysis of AI and linguistic diversity draws attention to how dominant-language technologies can accelerate convergence in smaller language communities without deliberate policy intent [33]. He cautions that institutional adoption of language technologies amplifies their influence, particularly when they interact with already vulnerable linguistic ecosystems.

Taken together, these historical and comparative cases caution against viewing language policy as merely technical or instrumental. They support the argument that when municipalities adopt tools that systematically reshape official language, including AI-driven simplification systems, they participate in shaping the broader linguistic environment. Decisions about AI use in municipal communication therefore function as choices with long-term societal implications, not merely efficiency measures.

4 Positions and Scenarios

This section outlines four policy positions regarding the use of AI in municipal text simplification. It distinguishes between approaches that delegate authorship to AI, approaches that retain AI as a co-author under human oversight, and approaches that restrict AI to diagnostic and learning-oriented support. A final position rejects AI use altogether. The section compares these scenarios to clarify their underlying assumptions, benefits, and risks, and to situate the paper’s defended position within the broader policy landscape.

4.1 Full AI-driven rewriting and generation

The first position supports the full adoption of AI systems as active authors of municipal communication. In this scenario, AI tools rewrite official texts to B1 level and generate final versions that are sent to residents with limited or no human revision. The core rationale is efficiency combined with accessibility. Evidence of declining literacy reported by the Algemene Rekenkamer and by international organisations such as the OECD suggests that a substantial share of the population struggles with reading formal administrative texts [1, 27]. National oversight bodies, including the Inspectie van het Onderwijs, similarly document persistent pressure on reading competence across educational tracks [19, 20]. From this perspective, relying on AI to simplify communication appears as a proportionate response to a systemic problem.

Supporters of this position further point to empirical evidence that even carefully produced human texts often fail to meet comprehension goals. Corsius and colleagues show in the Monitor Begrijpelijkheid overheidsteksten that misunderstandings remain widespread across policy domains despite longstanding commitments to plain language [10]. AI-based rewriting is therefore framed not as a radical departure, but as a scalable extension of existing accessibility efforts.

The main benefit of full AI-driven rewriting is immediate and measurable improvement in surface-level readability, combined with reduced workload and faster production cycles. This position aligns closely with national

guidance that promotes B1 as the default standard for public communication, as articulated by CommunicatieRijk and related government bodies [9, 32]. Within this framework, AI functions as a neutral efficiency tool that helps municipalities comply with established policy expectations.

The central weakness of this position lies in the structural effects of automation. Critiques of the CEFR framework by Weir and Morales indicate that B1 is an imprecise and context-insensitive benchmark, particularly when applied outside formal language assessment settings [25, 38]. Embedding this benchmark into an automated workflow risks standardising an already blunt instrument. Research on AI-generated language further suggests that routine AI authorship promotes convergence toward standardised phrasing and reduced lexical diversity, a pattern documented by Fredrick and Craven and critically analysed by Bauer [5, 15]. Full AI authorship therefore risks prioritising efficiency at the expense of linguistic nuance and institutional judgment.

4.2 AI assisted rewriting with human approval

The second position endorses AI assisted rewriting while retaining formal human approval of final texts. In this model, AI systems generate simplified drafts or alternative formulations, which municipal staff then review, edit, and approve. The rationale is that this approach captures efficiency gains while maintaining accountability. AI remains an active co-author, but humans serve as gatekeepers.

The appeal of this position lies in its apparent balance. It acknowledges accessibility pressures and limited resources while preserving formal responsibility with municipal employees. Empirical research on text simplification suggests that AI assisted rewriting can improve comprehension and reduce cognitive load for readers, lending support to this pragmatic compromise [18].

However, this position carries a subtler risk. Studies of AI assisted writing indicate that even when humans remain in the loop, repeated exposure to AI-generated phrasing can influence writing habits and norms. Fredrick and Craven show that writers tend to converge toward machine-suggested patterns over time, even when they retain editorial control [15]. Bauer similarly argues that AI output normalises dominant lexical and syntactic forms, making human approval increasingly procedural rather than substantive [5]. In institutional contexts marked by workload pressure and efficiency expectations, this dynamic risks shifting responsibility from active authorship to passive validation.

4.3 Diagnostic and learning-oriented AI support

The third position, which this paper defends, advocates the use of AI as a diagnostic and educational instrument rather than as an author. In this scenario, AI tools analyse texts to flag complexity, identify abstract formulations, and explain why certain passages may be difficult for specific audiences. AI does not produce final versions of official texts and does not autonomously rewrite them.

The core rationale of this position is institutional capacity building. Research on government communication shows that misunderstandings often arise from legal framing and institutional logic rather than vocabulary alone. Corsius and colleagues demonstrate that accessibility problems frequently stem from how information is structured and contextualised [10]. Diagnostic AI use can make these issues visible and support staff learning without flattening expression.

This approach also directly addresses concerns about linguistic homogenization. By positioning AI as an analytical aid rather than a default author, municipalities avoid allowing probabilistic models to silently set linguistic norms. This aligns with critical research on AI-generated language by Sourati and colleagues, who emphasise the importance of preserving diversity and reflective judgment in institutional contexts [35]. The benefit is long-term resilience rather than short-term efficiency.

The primary weakness of this position is that it demands investment in training and organisational discipline. It does not maximise immediate efficiency gains and requires explicit governance choices. However, historical

research on institutional language shaping shows that cautious, reflective approaches are justified when public decisions influence language over time [8]. This position therefore represents a deliberate policy choice rather than a compromise.

4.4 Rejection of AI for this use case

The fourth position rejects the use of AI in municipal text simplification altogether. Its strongest motivation is the protection of language as a public good. From this perspective, civic language should retain richness, precision, and the capacity to express complex rights and obligations. Language is treated not as an efficiency problem but as a core component of democratic participation. This view aligns with concerns raised in international policy discourse, including UNESCO-linked analyses warning about the societal risks of linguistic flattening [30].

The ethical appeal of this position is substantial. It prioritises democratic depth and resists technological solutions to structural educational problems. It avoids delegating communicative responsibility to systems trained on past language patterns rather than civic intent.

Despite these strengths, the position is difficult to sustain institutionally. Evidence on literacy decline and persistent communication failures shows that existing human-only processes already struggle to meet accessibility demands, as documented by the Algemene Rekenkamer and by studies of government communication [1, 10]. Given widespread informal AI use, full rejection may also lead to ungoverned adoption rather than meaningful control. For these reasons, while ethically coherent, this position is unlikely to provide a workable policy framework for municipalities such as Oostzaan.

5 Rebuttals and Counterarguments

Several counterarguments are commonly advanced in favour of AI-based simplification in municipal communication. This section engages explicitly with these claims and explains why, despite their intuitive appeal and partial empirical support, they do not justify treating automated or semi-automated simplification as a safe default for public-sector communication.

5.1 Accessibility urgency and the limits of surface simplification

A first strong counterargument holds that accessibility needs are urgent and ethically compelling. Proponents argue that when large numbers of citizens struggle with reading, municipalities have a responsibility to prioritise immediate comprehensibility over concerns about linguistic richness or long-term effects. This claim is supported by evidence from the Algemene Rekenkamer, which documents persistent low literacy levels among adults, and by the Inspectie van het Onderwijs, which reports sustained pressure on reading competence across educational tracks [1, 19]. From this perspective, AI assisted B1 rewriting is framed as a pragmatic response to an existing exclusion problem rather than a speculative risk.

This argument is compelling as it recognises that current communication practices already fail many residents. However, research on government texts suggests that simplification alone does not reliably address the sources of misunderstanding. Corsius and colleagues show that comprehension problems frequently arise from legal framing, institutional logic, and implicit assumptions rather than from surface linguistic complexity [10]. Automated rewriting therefore risks standardising texts without improving conceptual clarity. The rebuttal does not deny the urgency of accessibility, but questions whether AI-based simplification meaningfully resolves the underlying problem or merely reformats it.

5.2 Human-in-the-loop models as a presumed safeguard

A second counterargument claims that the risks of AI-based simplification are overstated, particularly when AI systems are used under human supervision. Supporters emphasise that large language models are tools rather

than autonomous actors and argue that their societal impact depends on patterns of use. They further point to emerging evidence that AI assisted simplification can improve comprehension and reduce cognitive load for readers. Experimental work summarised by Ardila and Kakarmath and large-scale evaluations conducted by Guidroz and colleagues provide empirical support for these short-term benefits [3, 18].

This counterargument is credible, and the evidence base on benefits is still developing. However, research on AI assisted writing complicates the assumption that human oversight reliably neutralises longer-term effects. Fredrick and Craven demonstrate that repeated exposure to AI-generated phrasing leads writers to converge toward machine-suggested patterns, even when they retain formal editorial control [15]. Bauer similarly argues that AI output normalises dominant lexical and syntactic forms, contributing to reduced linguistic diversity over time [5]. In institutional settings characterised by workload pressure and routinised review, human approval risks becoming procedural rather than substantive.

5.3 AI as a neutral efficiency aid within existing B1 policy

A third counterargument holds that AI simplification merely formalises what municipalities already aim to do manually. Since B1 is the endorsed standard and staff are trained accordingly, AI tools are framed as neutral efficiency aids rather than as a qualitative shift in language policy. This argument accurately reflects current guidance promoted by CommunicatieRijk and related policy bodies, which present B1 as a pragmatic default for public communication [9, 32].

However, critiques of the CEFR framework complicate this claim. Weir argues that CEFR levels abstract away from contextual and cognitive dimensions of comprehension, making them blunt instruments outside language testing contexts [38]. Morales provides empirical evidence that grammatical and structural complexity does not align cleanly with CEFR bands, undermining the assumption that texts can be reliably calibrated to a single level such as B1 [25]. Siripol and colleagues further demonstrate that automated CEFR classification tools assign inconsistent levels to identical texts, raising questions about their suitability for routine policy enforcement [34]. Automating an already imprecise standard risks amplifying its limitations rather than resolving them.

5.4 Inevitability and the case for conditional adoption

A final counterargument holds that limiting AI use is unrealistic given its widespread societal adoption. If citizens and municipal employees already use AI tools informally, formal integration is presented as a way to maintain transparency and control. This pragmatic reasoning supports the view that AI adoption cannot be fully prevented.

However, research on institutional language shaping suggests that how tools are adopted matters as much as whether they are adopted. Comparative studies by Bokhorst-Heng and Silver and by Gonzalez show that institutional language choices can establish and reinforce specific norms and hierarchies over time, even when adoption is framed as neutral or inevitable [8, 17]. Conditional and diagnostic use recognises inevitability while still exercising responsibility over linguistic impact.

Taken together, these rebuttals acknowledge that AI assisted simplification responds to real pressures and offers measurable short-term benefits. The position defended here does not deny those benefits. It argues instead that full reliance on automated or semi-automated simplification is a disproportionate response given institutional constraints, unresolved uncertainties, and the risk of long-term linguistic homogenization.

6 Recommendations for Practice

For municipal decision makers, the position defended in this paper requires translating observed effects from the existing literature into concrete governance boundaries. Rather than treating AI as a neutral productivity tool, municipalities should articulate explicit policy choices about which applications are acceptable, which are

discouraged, and which are prohibited. This approach reflects how public institutions typically manage risk and responsibility in other sensitive domains.

Acceptable uses of AI include diagnostic analysis of texts, explanation of linguistic complexity, identification of jargon, and educational support for staff learning. Discouraged uses include routine AI-generated rewriting that staff approve without substantive reflection. Prohibited uses include autonomous AI authorship of legally binding or norm-setting municipal communication. These distinctions are not technical constraints but policy choices about where responsibility for civic language should reside.

For municipal practitioners and policy advisors, the central recommendation is therefore not to reject AI outright, but to avoid positioning it as an autonomous writer that routinely rewrites official texts to B1 level. Instead, AI should be deployed as a diagnostic and supportive instrument that strengthens human judgment. Research on government communication by Corsius and colleagues shows that misunderstandings often arise from legal framing and institutional logic rather than vocabulary alone [10]. AI tools can be configured to highlight such sources of complexity by flagging dense constructions, abstract references, or procedural assumptions, enabling staff to make informed revisions.

Second, outputs produced by AI should function as signals for critical review rather than just as candidate replacements. When AI tools flag passages as complex or problematic, these signals should prompt trained staff to analyse why a formulation may hinder comprehension and to decide whether and how revision is warranted. The task is not to select among AI-generated alternatives, but to exercise informed judgment over the underlying communicative issue. This approach directly addresses critiques of CEFR-based simplification. As Weir and Morales have shown, B1 is an imprecise and context-insensitive benchmark when applied mechanically [25, 38]. Treating AI flags as diagnostic cues rather than default solutions helps prevent the uncritical adoption of standardised phrasing and preserves flexibility in tone, register, and legal precision.

Third, municipalities should invest in linguistic competence alongside technological tools. Training employees to understand why texts are difficult and how language choices affect comprehension reduces overreliance on automation. Empirical studies of AI assisted writing by Fredrick and Craven and by Bauer indicate that repeated exposure to homogenized AI output can influence writing habits and narrow expressive range [5, 15]. Active reflection and professional development can counterbalance this effect and sustain institutional language capacity.

Fourth, AI use should be embedded within an explicit language policy. Municipalities should define clear boundaries for acceptable AI applications, distinguishing between diagnostic support, drafting assistance, and prohibited autonomous rewriting for legally binding texts. Periodic review of AI-supported communication can help detect unintended standardisation or loss of nuance over time, a concern raised in the literature on linguistic homogenization and institutional language shaping [8, 35].

Finally, municipalities should treat AI as part of a broader accessibility strategy rather than as a standalone solution. Simplification should be combined with clearer information structuring, better explanation of procedures, and alternative communication channels where appropriate. Guidance collected by the Kenniscentrum Begrijpelijke Taal and CommunicatieRijk consistently emphasises that accessibility is an institutional responsibility, not merely a linguistic one [9, 22]. Aligning AI use with this broader strategy helps ensure that accessibility gains do not come at the expense of long-term linguistic resilience.

7 Discussion and Limitations

The position advanced in this paper is explicitly concerned with how municipalities ought to act under conditions of uncertainty. It does not claim to establish deterministic causal effects between AI assisted simplification and linguistic decline. Rather, it offers an argued judgment about how municipalities should act under conditions of uncertainty, given their public responsibility for civic communication.

The argument rests on what is currently known and can be inferred based on existing literature from different domains. Research on declining literacy in the Netherlands is robust and convergent. Reports by the Algemene Rekenkamer and international assessments coordinated by the OECD consistently show that reading competence is under pressure across age groups and educational tracks [1, 27]. National oversight bodies, including the Inspectie van het Onderwijs, reinforce this picture by documenting persistent shortcomings in basic language skills and widening disparities between schools [19, 20]. Studies of government communication likewise demonstrate enduring accessibility failures that justify institutional intervention, as shown by Corsius and colleagues and synthesised in policy guidance collected by the Kenniscentrum Begrijpelijke Taal [10, 22].

Evidence on the long-term linguistic effects of AI-generated language remains more limited, but it is directionally consistent. Empirical studies by Fredrick and Craven and by Bauer indicate convergence toward dominant lexical and syntactic patterns in AI assisted writing [5, 15]. Sourati and colleagues extend this analysis by arguing that widespread reliance on shared language models may reduce cognitive and linguistic diversity at scale [35]. While much of this research focuses on educational or experimental contexts, historical and comparative studies of institutional language policy suggest that even modest, systematic shifts can accumulate over time when embedded in public systems [8, 17].

A central limitation of the argument is uncertainty. AI systems evolve rapidly, and institutional usage patterns are still stabilising. Some risks discussed here may prove smaller than anticipated or may be mitigated through future design and governance choices. However, the absence of definitive long-term evidence does not imply the absence of risk. Research across domains consistently shows that language technologies are not neutral and that institutional adoption amplifies their influence, particularly when combined with existing pressures on literacy and accessibility.

Municipal decision making does not require certainty to act responsibly. What matters is proportionality and reversibility. Automated rewriting embedded in municipal workflows is difficult to reverse once normalised, whereas diagnostic and conditional use preserves flexibility and human judgment. Given the public role of municipalities in shaping civic language, uncertainty therefore justifies restraint rather than maximal automation.

The position defended here treats uncertainty not as a reason for inaction, but as a reason to limit AI authorship while enabling targeted support. By restricting AI to diagnostic and learning-oriented functions, municipalities can respond to accessibility challenges while retaining the capacity to revise their approach as linguistic, technological, and empirical conditions evolve.

8 Conclusion

Dutch municipalities such as Oostzaan face genuine and pressing pressure to improve the accessibility of official communication. Evidence of declining literacy and persistent misunderstandings in government texts makes inaction untenable. The Algemene Rekenkamer documents the scale of low literacy among adults, while studies of government communication show that even well-intentioned human texts frequently fail to achieve clarity [1, 10]. These findings justify active institutional efforts to improve how municipalities communicate with residents.

At the same time, the rapid availability of AI writing and rewriting tools introduces a policy choice that is not merely technical, but institutional and tied to questions of responsibility. Research on the limitations of CEFR-based simplification shows that B1 is an imprecise proxy for real-world comprehension, particularly when applied mechanically outside language assessment contexts [25, 38]. Evidence on AI-generated language further indicates that routine reliance on automated rewriting promotes standardised phrasing and reduced expressive diversity, with potential long-term effects on institutional language practices [5, 35].

This paper has argued that structural reliance on AI-driven B1 rewriting, including models that rely on routine human approval, is therefore an insufficient and potentially risky response. Automated simplification may improve surface readability, but it also risks narrowing linguistic expression, entrenching standardised administrative

registers, and shifting responsibility for communicative clarity away from human judgment. Once embedded in workflows, such practices are difficult to reverse and may reshape civic language in ways that are neither transparent nor deliberate.

Municipal responsibility extends beyond efficiency and guideline compliance. Public institutions shape how rights, obligations, and democratic participation are articulated. For this reason, the defensible policy choice is not maximal automation, but governed restraint. By prohibiting AI from acting as an autonomous author while actively deploying it for diagnostic analysis, explanation of complexity, and staff learning, municipalities can improve accessibility without surrendering linguistic judgment.

The primary risk for municipalities is not only the misuse of AI, but the failure to set boundaries early. In the absence of explicit policy choices, informal and unreflective adoption may define practice by default. Clear governance decisions, limits on AI authorship, and ongoing evaluation are therefore essential. Responsible integration of AI into municipal communication requires aligning innovation with long-term linguistic resilience and democratic accountability rather than treating accessibility as a problem to be automated away.

9 Citations and Bibliographies

References

- [1] Algemene Rekenkamer. 2016. Aanpak van laaggeletterdheid. <https://www.rekenkamer.nl/documenten/2016/04/20/aanpak-van-laaggeletterdheid> Rapport 20-04-2016.
- [2] Barrett R. Anderson, Jash Hemant Shah, and Max Kreminsk. 2024. Creativity Support Tools and Idea Homogenization. *Commun. ACM* , (2024), . <https://dl.acm.org/doi/fullHtml/10.1145/3635636.3656204>
- [3] Diego Ardila and Sujay Kakarmath. 2025. Making complex text understandable: Minimally-lossy text simplification with Gemini. <https://research.google/blog/making-complex-text-understandable-minimally-lossy-text-simplification-with-gemini/> Google Research Blog.
- [4] Placeholder Author. 2025. Decolonizing English Language Teaching in the Philippines. *Journal of Language, Diversity and Learning* , (2025), . <https://journal.formosapublisher.org/index.php/jldl/article/download/14720/14587/64002>
- [5] Nicolas Bauer. 2025. Does ChatGPT Increase Language Homogenization? , , 11–31. doi:10.1007/978-3-658-46344-1_2
- [6] Belastingdienst. n.d.. Begrijpelijke brieven als startpunt voor heldere communicatie. <https://over-ons.belastingdienst.nl/begrijpelijke-brieven-als-startpunt-voor-heldere-communicatie/> Accessed December 2025.
- [7] Belastingdienst Toeslagen. n.d.. Communicatie in eenvoudige taal. <https://www.overtoeslagen.nl/onderwerpen/v/verbeteren-dienstverlening/communicatie-in-eenvoudige-taal> Accessed December 2025.
- [8] Wendy D. Bokhorst-Heng and Rita Elaine Silver. 2017. Reconsidering language shift within Singapore’s Chinese community. *International Journal of the Sociology of Language* 2017, 243 (2017), 135–170. doi:10.1515/ijsl-2017-0037
- [9] CommunicatieRijk. n.d.. Inclusieve communicatie. <https://www.communicatierijk.nl/vakkennis/uitgangspunten-en-organisatie/inclusieve-communicatie/inclusieve-communicatie1> Accessed December 2025.
- [10] Mischa Corsius, Vera Lange, Yvette Linders, Henk Pander Maat, Els van der Pool, Nina Sangers, Keun Sliedrecht, Wouter Sluis-Thiescheffer, and Charlotte Swart. 2023. *Monitor Begrijpelijkheid overheids teksten 2022*. Technical Report. HAN University of Applied Sciences and Universiteit Utrecht. <https://taalunie.org/publicaties/218/monitor-begrijpelijkheid-overheids-teksten-2022>
- [11] Council of Europe. n.d.. The CEFR Levels. Online. <https://www.coe.int/en/web/common-european-framework-reference-languages/level-descriptions> Accessed 26 November 2025.
- [12] Door Lotje. 2025. Bereik jouw doelgroep met teksten in taalniveau B1. <https://doorlotje.nl/blog/laaggeletterden> Citing Stichting Lezen & Schrijven estimates.
- [13] European Commission. 2025. Education and Training Monitor: Country Report Netherlands. <https://op.europa.eu/webpub/eac/education-and-training-monitor/en/country-reports/netherlands.html> Accessed December 2025.
- [14] Hind Farhan. 2025. The Impact of AI-Powered Writing Tools on Students’ Writing Performance. doi:10.13140/RG.2.2.22893.09445
- [15] Daniel R. Fredrick and Laurence Craven. 2025. Lexical diversity, syntactic complexity, and readability. *Frontiers in Education* 10 (2025), 1616935. doi:10.3389/feduc.2025.1616935
- [16] Jad M. Gayed et al. 2022. Exploring an AI-based writing assistant’s impact on English as a foreign language writing. *Journal of Computer Assisted Learning* , (2022), . <https://www.sciencedirect.com/science/article/pii/S2666920X22000108>
- [17] Andrew Gonzalez. 1998. The language planning situation in the Philippines. *Journal of Multilingual and Multicultural Development* 19, 5 (1998), 487–525. doi:10.1080/01434639808666365
- [18] Theo Guidroz, Diego Ardila, Jimmy Li, Adam Mansour, et al. 2025. LLM-based Text Simplification and its Effect on User Comprehension and Cognitive Load. *arXiv preprint* , (2025), . arXiv:2505.01980 [cs.CL] doi:10.48550/arXiv.2505.01980
- [19] Inspectie van het Onderwijs. 2021. *The State of Education in the Netherlands 2021*. Technical Report. Inspectorate of Education, Netherlands. <https://english.onderwijsinspectie.nl/documents/annual-reports/2021/04/30/the-state-of-education-in-the-netherlands-2021>
- [20] Inspectie van het Onderwijs. 2024. *The State of Education 2024*. Technical Report. Inspectorate of Education, Netherlands. <https://english.onderwijsinspectie.nl/documents/publications/2024/08/7/the-state-of-education>
- [21] Inspectie van het Onderwijs. 2025. *The State of Education 2025*. Technical Report. Inspectorate of Education, Netherlands. <https://english.onderwijsinspectie.nl/documents/publications/2025/09/23/the-state-of-education-2025> Accessed December 2025.
- [22] Kenniscentrum Begrijpelijke Taal. n.d.. Duidelijke Taal. <https://www.kcbr.nl/themas/themas-z/duidelijke-taal> Accessed December 2025.
- [23] MIT Media Lab. 2024. Your Brain on ChatGPT. <https://www.media.mit.edu/publications/your-brain-on-chatgpt/> Accessed December 2025.
- [24] Kaijie Mo and Renfen Hu. 2024. ExpertEase: A Multi-Agent Framework for Grade-Specific Document Simplification with Large Language Models. In *Findings of the Association for Computational Linguistics: EMNLP 2024*. Association for Computational Linguistics, Miami, Florida, USA, 9080–9099. doi:10.18653/v1/2024.findings-emnlp.530
- [25] Gabriel Morales. 2021. *A text grammar-based index for text complexity across CEFR levels*. Ph. D. Dissertation. Universidad Metropolitana de Ciencias de la Educación. https://www.researchgate.net/publication/362113418_A_text_grammar-based_index_for_text_complexity_across_CEFR_levels

- [26] C. L. P. Ng. 2014. Mother tongue education in Singapore: concerns, issues and controversies. *Current Issues in Language Planning* 15 (10 2014), 361–375. doi:10.1080/14664208.2014.927093
- [27] OECD. 2019. PISA 2018 Results (Volume I): What Students Know and Can Do. https://www.oecd.org/en/publications/2019/12/pisa-2018-results-volume-i_947e3529.html OECD Programme for International Student Assessment.
- [28] OECD. 2025. Education at a Glance 2025: Netherlands. https://www.oecd.org/en/publications/education-at-a-glance-2025_1a3543e2-en/netherlands_26a700bc-en.html Accessed December 2025.
- [29] OECD and NL Times. 2023. Dutch kids' reading, maths and science skills declining: OECD. <https://nltimes.nl/2023/12/05/dutch-kids-reading-maths-science-skills-declining-oecd> Summary of PISA 2022 results for the Netherlands.
- [30] Aristotelis Ioannis Paschalidis. 2025. AI and the Great Linguistic Flattening. *UNESCO*, (06 2025), . <https://www.unesco.org/en/articles/ai-and-great-linguistic-flattening> Last updated 27 October 2025.
- [31] Pharos. n.d.. Begrijpelijke communicatie is goed voor iedereen. <https://www.pharos.nl/infosheets/begrijpelijke-communicatie-is-goed-voor-iedereen/> Accessed November 2025.
- [32] Rijksoverheid. n.d.. Begrijpelijke communicatie. <https://www.communicatierijk.nl/vakkennis/uitgangspunten-en-organisatie/inclusieve-communicatie/begrijpelijke-communicatie> Accessed November 2025.
- [33] Gerald Roche. 2024. Four Villages in Tibet Have a Lot to Tell Us About Artificial Intelligence and the Future of Linguistic Diversity. <https://www.cornellpress.cornell.edu/four-villages-in-tibet-and-the-limits-of-ai/> Cornell University Press authors' blog.
- [34] Piyathat Siripol, Seongha Rhee, Suthathip Thirakunkovit, and Aphiwit Liang-Itsara. 2025. Evaluating the consistency of automated CEFR analyzers: a study of English language text classification. *International Journal of Evaluation and Research in Education* 14, 4 (2025). doi:10.11591/ijere.v14i4.33528
- [35] Zhivar Sourati, Alireza S. Ziabari, and Morteza Dehghani. 2025. The Homogenizing Effect of Large Language Models on Human Expression and Thought. *arXiv preprint abs/2508.01491* (2025), . doi:10.48550/arXiv.2508.01491
- [36] Stichting Lezen & Schrijven. n.d.. Eenvoudige taal. <https://www.lezenenschrijven.nl/wat-doen-wij/oplossing-voor-je-vraagstuk/eenvoudige-taal> Accessed November 2025.
- [37] Anne L. Verdegaal. 2021. The Dutch decline in PISA reading performance explained. , (2021), . <https://essay.utwente.nl/86741/>
- [38] Cyril J. Weir. 2005. Limitations of the Common European Framework for developing comparable examinations and tests. *Language Testing* 22, 3 (2005), 281–300. doi:10.1191/0265532205lt309oa